

## HW 1: Parallel Programming on a Multicore Multiprocessor

### Part 1. Shared-Memory Programming with Threads

Compile and execute the program in the file `compute_pi.c`, which computes an estimate of  $\pi$  using the parallel algorithm discussed in class. It should be compiled and executed on `grace.hprc.tamu.edu`.

Load the Intel software stack prior to compiling and executing the code.

```
module load intel
```

To compile, use the command:

```
icx -o compute_pi.exe compute_pi.c -lpthread
```

To execute the program, use

```
./compute_pi.exe <n> <p>
```

where `<n>` represents the number of points and `<p>` represents the number of threads. The output of a sample run is shown below.

```
./compute_pi.exe 1000000 4
```

```
Trials = 1000000, Threads = 4, pi = 3.1433480000, error = 5.59e-04,  
time (sec) = 0.0043
```

The run time of the code should be measured when it is executed in dedicated mode. Use the batch file `compute_pi.grace_job` to execute the code in dedicated mode using the following command on Grace:

```
sbatch compute_pi.grace_job
```

Hint: Copy-paste from this document might not work because of “ghost” characters that get added. Type the entire command directly into the terminal shell on the HPRC computer.

1. Execute the code for  $n=10^8$  with  $p$  chosen to be  $2^k$ , for  $k = 0, 1, \dots, 13$ . Using the experimental data obtained from these experiments, answer the following questions. For plots, use a logarithmic scale for the x-axis.
  - 1.1. (10 points) Plot execution time versus  $p$  to demonstrate how time varies with the number of threads.
  - 1.2. (10 points) Plot speedup versus  $p$  to demonstrate the change in speedup with  $p$ .
  - 1.3. (5 points) Using the definition:  $\text{efficiency} = \text{speedup}/p$ , plot efficiency versus  $p$  to demonstrate how efficiency changes as the number of threads are increased.
  - 1.4. (5 points) In your experiments, what value of  $p$  minimizes the parallel runtime?
2. Repeat the experiments with  $n=10^{10}$  to obtain the execution time for  $p=2^k$ , for  $k = 0, 1, \dots, 13$ .
  - 2.1. (5 points) In this case, what value of  $p$  minimizes the parallel runtime?
  - 2.2. (5 points) Do you expect the runtime to increase as  $p$  is increased beyond a certain value? If so, why? And is this observed in your experiments.
3. (5 points) Do you expect that there would be a difference in the number of threads needed to obtain the minimum execution time for two values of  $n$ ? Is this observed in your experiments.

4. (5 points) Plot error versus  $n$  to illustrate accuracy of the algorithm as a function of  $n$ . You may have to run experiments with different values of  $n$ ; for example  $n$  could be chosen to be  $10^k$ , for  $k = 3, \dots, 9$ . Use  $p = 48$ .

## Part 2. Distributed-Memory Programming with MPI

Compile and execute the program in the file `compute_pi_mpi.c`, which computes an estimate of  $\pi$  using the parallel algorithm discussed in class. It should be compiled and executed `grace.hprc.tamu.edu`.

Load the Intel software stack prior to compiling and executing the code.

```
module load intel
```

To compile, use the command:

```
mpiicx -o compute_pi_mpi.exe compute_pi_mpi.c
```

To execute the program, use

```
mpirun -np <p> ./compute_pi_mpi.exe <n>
```

where  $\langle n \rangle$  represents the number of intervals and  $\langle p \rangle$  represents the number of processes. The output of a sample run is shown below.

```
mpirun -np 4 ./compute_pi_mpi.exe 100000000
```

```
n = 100000000, p = 4, pi = 3.1415926535897749, relative error = 5.80e-15, time (sec) = 0.0608
```

The run time of the code should be measured when it is executed in dedicated mode. Use the batch file `compute_pi_mpi.grace_job`, to execute the code in dedicated mode using the following command on Grace:

```
sbatch compute_pi_mpi.grace_job
```

5. Execute the code for  $n=10^8$  with  $p$  chosen to be  $2^k$ , for  $k = 0, 1, \dots, 6$ . Specify `ntasks-per-node=4` in the job file. Using the experimental data obtained from these experiments, answer the following questions. For plots, use a logarithmic scale for the x-axis.
  - 5.1. (10 points) Plot execution time versus  $p$  to demonstrate how time varies with the number of processes.
  - 5.2. (10 points) Plot speedup versus  $p$  to demonstrate the change in speedup with  $p$ .
  - 5.3. (5 points) Using the definition:  $\text{efficiency} = \text{speedup}/p$ , plot efficiency versus  $p$  to demonstrate how efficiency changes as the number of processes is increased.
  - 5.4. (5 points) What value of  $p$  minimizes the parallel runtime?
6. (10 points) With  $n=10^{10}$  and  $p=64$ , determine the value of `ntasks-per-node` that minimizes the `total_time`. Plot time versus `ntasks-per-node` to illustrate your experimental results for this question.
7. Execute the code with  $p=64$  for  $n=10^2, 10^4, 10^6$  and  $10^8$ , with `ntasks-per-node=4`.
  - 7.1. (5 points) Plot the speedup observed as a function of  $n$  on  $p=64$  w.r.t.  $p=1$ . You will need to obtain execution time on  $p=1$  for  $n=10^2, 10^4, 10^6$  and  $10^8$ .

## CSCE 735 Summer 2025

- 7.2. (5 points) Plot the relative error versus  $n$  to illustrate the accuracy of the algorithm as a function of  $n$ .

**Submission:** Upload a single PDF or MSWord document with your answers to Canvas.